Alex Fernsby

Copenhagen, Denmark
M +1 555 0100 • **E** alex.fernsby@example.com • **gh** afernsby-example

Profile

Data engineer with 7 years of experience building batch and streaming pipelines, most recently for logistics and earlier for retail. Strongest in Python, SQL and Apache Airflow; I own shipment-event processing end to end, work closely with the operations teams who use the data, and mentor junior engineers.

Core Competencies

- **Python and SQL**: ingestion services, reporting and warehouse development.
- **Pipelines**: Apache Airflow, Kafka and dbt for batch and streaming workloads.
- **Data platforms**: PostgreSQL and BigQuery; strong Docker skills.
- **Data quality**: Great Expectations, data contracts, dbt data tests and row-by-row validation during warehouse migrations.
- **Delivery**: working knowledge of GitHub Actions and Terraform for deployments; Looker for operations reporting.

Professional Experience

Tidewater Logistics — **Copenhagen**
Senior Data Engineer — *March 2022 - present*

- Own the shipment-events pipeline handling 40 million events daily from Kafka into PostgreSQL and BigQuery, orchestrated with Apache Airflow.
- Cut the delivery-performance mart's nightly rebuild from 6 hours to 55 minutes using incremental dbt models.
- Mentor two junior engineers and lead the weekly data-modelling review.
- Introduced Great Expectations tests on 30 core tables; failed loads dropped 70% in the first quarter.
- Moved deployments to GitHub Actions and Terraform; new environments take 20 minutes instead of two days.
- Defined the shipment-events data contract with five producing teams and wrote the on-call runbook for the shipment-events pipeline.
- Cut the BigQuery events warehouse bill by 35% through partitioning and clustering the five largest tables.
- Built a Looker explore for the operations team so planners answer delivery-delay questions themselves instead of filing tickets.
- Ran interviews for three data engineering hires and wrote the team's take-home exercise.

Marlow Retail Group — **Utrecht**
Data Engineer — *June 2018 - February 2022*

- Built the company's first Airflow deployment and migrated 45 cron jobs onto it.
- Designed a PostgreSQL customer-orders star schema used by 12 analysts.
- Wrote Python ingestion services for REST and SFTP partner price feeds, handling about 2 million rows daily.
- Supported the move of the warehouse from on-premise PostgreSQL to BigQuery, validating 60 tables row by row.
- Trained 12 analysts on dbt and reviewed their first 40 models.
- Added alerts to 45 Airflow jobs; late feeds are noticed within 15 minutes instead of the next morning.
- Wrote the migration guide for moving reports off the old warehouse, used by four teams.

Delft Energy Cooperative **Delft**
Junior Analyst *September 2016 - May 2018*

- Automated monthly consumption reports with Python and SQL, saving the team about 3 days each month.
- Built a small Flask dashboard that 25 cooperative members used to read their own consumption.
- Cleaned five years of meter readings with Pandas and documented the 14 data-quality rules that came out of it.
- Presented the consumption findings to the cooperative's board twice a year.

Education

Delft University of Technology
MSc Computer Science (data systems) *2016*

Delft University of Technology
BSc Computer Science *2014*

Projects

Open-source dbt package for freight-rate data tests; 120 GitHub stars. freight-tests

2024: Talk: "Data contracts in practice", Rotterdam Data Meetup.

Certifications

2023: Google Cloud Professional Data Engineer

2022: dbt Analytics Engineering Certification

Languages

English: C1

Dutch: Native

German: A2

Publications

2024: "Testing freight data with dbt", Data Engineering Weekly guest post.

Honors and Awards

2023: Tidewater Logistics Engineering Excellence Award

Work Authorisation and Availability

Work rights: Dutch citizen; authorised to work anywhere in the EU without sponsorship.

Availability: Three months' notice; earliest start 1 February 2027. Full-time, at most 40 hours per week.